


# Estimator complexity in Deep Reinforcement Learning



UNIVERSITY OF  
BUCHAREST  
VIRTUTE ET SAPIENTIA

Carla Crivoi, Andrei Hodoroaga and Florin Radu Gogianu \*

University of Bucharest, Romania

\*Bitdefender, Romania

crivoicarla02@gmail.com, andrei.hodoroaga36@gmail.com, florin.gogianu@bitdefender.com

## 1. Introduction

The project's goal is to determine the effect of different objective functions used in RL on the complexity of the neural networks they are trained on.

### Challenges:

1. isolating the possible effects of each algorithm on the neural network
2. determining the proper metrics to evaluate the network on

## 2. Dataset and Preprocessing

### Dataset Description:

The dataset contains 10,000 samples from a **Grid World** environment, a rectangular grid with a start and end state, and (optionally) some walls in between.

The samples come from complete trajectories of an agent following a random policy (having equal probability to move to every possible next state).

One sample of the dataset is a tuple of the form  $(\{x, y\}, z)$ , where  $\{x, y\}$  are the coordinates on the grid and  $z$  is the associated label.  $z$  differs between the algorithms:

- for the regression task it is the value function for the state  $\{x, y\}$
- for Monte Carlo and TD(0) it is the respective objective function at the state  $\{x, y\}$




Figure 1: Example of a Grid World

**Data preprocessing:** The only preprocessing step implemented is normalizing the state coordinates between  $(-0.99, 0.99)$ .

## 3. Objectives

1. **Regression to the true value function** - We evaluate a random policy to convergence and then train the neural networks to approximate the true value function.
2. **Monte Carlo prediction** - We train the neural networks using the returns obtained from complete episodes for each state.
3. **Semi-gradient TD(0)** - We train the neural networks using the TD target:  $R_{t+1} + \gamma \hat{v}(S_{t+1}, w)$ .

## 4. Models


We trained the algorithms on several different architectures: with 1, 2, 3, and 4 hidden layers, and each layer with 2, 4 or 16 neurons.

## 5. Results


Training the models:

1. **Regression task and Monte Carlo:** We trained the networks until the loss was reduced to less than 0.0001 or until there was no improvement in the loss for 200 consecutive epochs.
2. **TD0 algorithm:** Each network was trained for around 5 epochs.

### Fourier complexity:




### Polynomial complexity:




### Visual representation of predictions:

#### V(s):




#### Regression:

#### TD(0):




#### Monte Carlo:




### Heatmaps of the neural networks for each algorithm:


#### Untrained:




#### Regression:



#### TD(0):



#### Monte Carlo:



## 6. Network complexity metrics

**Fourier complexity:** The spectral or frequency complexity of a function on a 2D surface is evaluated, using Fourier transforms. This is done by analyzing how the function values vary on a grid defined in a 2D space.

$$f(k) = \sum_{x \in X_{grid}} \omega_x^T k f(x)$$

**Polynomial complexity:** We evaluate how a function behaves over a 2D domain by approximating and analyzing it with Chebyshev polynomials. This can be useful in mathematical and numerical analysis to understand characteristics such as the degree of variation or smoothness of a function over a given domain.

$$CChebyshev(f) = \sum_{n_1, n_2=0}^N |c_{n_1 n_2}| \cdot \| [n_1, n_2] \|^2 \frac{1}{\sum_{n_1, n_2=0}^N |c_{n_1, n_2}|}$$

## 7. Limitations and future work

- Results based on only one initialization for each model
- Metrics limited to two types of complexity metrics (Chebyshev and Fourier), will add compressibility in the future
- TD(0) experiments should be revisited, we suspect under-training
- Only one MDP was used in the comparison, we should compare different MDP classes